

Week 1 Notes: Introduction to Probabilistic Modeling

In this collection of notes, we briefly outline the main thrust of this course by discussing (at a high level of generality) the topic of *probabilistic modeling* (or *statistical modeling*). Throughout this course we will focus on the application of probabilistic modeling with an eye towards solving concrete problems involving complex datasets. The following problems will recur throughout:

1. *How do we formalize the questions we would like to answer in a way that can be addressed using statistical modeling?*
2. *How do we design an appropriate probabilistic model that answers the questions we are interested in?*
3. *Computationally, how do we extract answers from a probabilistic model?*
4. *How do we properly account for uncertainty in our conclusions?*
5. *How do we check that our model is “good enough” as a representation of reality for our purposes, and ensure that the conclusions we draw are robust?*

1 What is a Probabilistic Model

At this stage in your education, you are undoubtedly familiar with the concept of probabilistic modeling. All this means is that:

1. We have a collection \mathcal{D} of data that we have measured.
2. We posit that \mathcal{D} has arisen by randomly sampling it according to some *data generating process* G , i.e., $\mathcal{D} \sim G$.
3. The data generating process G is assumed to lie in some collection \mathcal{G} of possible data generating processes.

That's it! The most generic probabilistic model stops after Step 2. Such a model is not very useful, however; for example, if $\mathcal{D} = (X_1, \dots, X_N)$ and all we know is that G is an arbitrary probability distribution on \mathbb{R}^N , then it will be more-or-less impossible to draw any conclusions about G from the data. Generally we assume more.

Example 1: Empirical Distribution

Suppose that $\mathcal{D} = (X_1, \dots, X_N)$ and suppose further that $X_i \stackrel{\text{iid}}{\sim} F$ so that $G = F^N$. Then a reasonable estimate of F is the *empirical distribution*

$$\mathbb{F}_N = \frac{1}{N} \sum_{i=1}^N \delta_{X_i},$$

where δ_{X_i} denotes the point mass distribution at X_i . We make no assumptions about F itself; it can be (say) a normal distribution, an exponential distribution, a Cauchy distribution, or any other distribution.

Part of the point of bringing up this example is to point that, while weak, the assumption that a collection of data is iid *is an assumption that should be scrutinized*. Data might fail to be iid if (for example) the X_i 's correspond to a time series.

2 Statistical Inferences

Prior to specifying a probabilistic model, I find it useful in many settings to begin by specifying some *concrete, model independent* inferences one would like to draw from the data, make some falsifiable predictions about what we expect the data to show, or specify some quantities of interest we would like to predict. Here are some reasons to do this:

1. It is useful as an informal step for guiding model specification, so that we do not waste energy modeling aspects of the data generating process that are not pertinent to our goals.
2. Paraphrasing Christian Robert's "Bayesian Choice", this helps us aim our analysis at an objective purpose (e.g., whether to send a specific drug to market, whether a robot should modify its behavior, etc), having measurable consequences. This can also provide objective benchmarks with which to assess the success of our models and estimation strategies.
3. When proceeding formally from a decision-theoretic perspective (which we mostly won't do), this provides benchmarks for comparing different estimation procedures.

Already in Example 1 interesting problems start occurring. For example, suppose we know that F has a density $f(x)$ that we are interested in estimating. How do we go about estimating $f(x)$ within the context of the iid model?

Example 2: Density Estimation

The **galaxy** dataset available in the **MASS** package contains “a numeric vector of velocities in km/sec of 82 galaxies from 6 well-separated conic sections of an unfilled survey of the Corona Borealis region. Multimodality in such surveys is evidence for voids and superclusters in the far universe.”

Our primary goal is to determine whether the density $f(x)$ of the velocities is multimodal. We’ll assume that the sample of galaxies is taken iid from some distribution $Y_i \stackrel{\text{iid}}{\sim} F$ with density $f(y)$. We then form the density estimate

$$\hat{f}(y) = \frac{1}{N} \sum_i \phi(y; Y_i, h)$$

where $\phi(y; \mu, \sigma)$ denotes the density of a $\text{Normal}(\mu, \sigma^2)$ distribution. This estimator is biased and, in fact, there are *no* unbiased estimators of $f(y)$ if we make no further assumptions on F .

A plot of the density estimate is given in Figure 1. Given our final aim of determining whether the density is multimodal, a complete statistical analysis might ask how accurate counting the number of modes of the kernel density estimator is as an estimator of the number of modes, and what the degree of uncertainty we have in the number of modes is, but to a first approximation it seems clear that the density is, in fact, multimodal.

In practice, models are chosen not just with the aim to adequately describe reality (a task for which they will always come up short) but also for their ability to answer specific questions about the data generating process. Statistical models are useful in that they allow us to both (i) attempt to answer our questions and (ii) attempt to determine how robust our answers by providing uncertainty quantification.

3 Very Simple Probabilistic Models

The estimator $\hat{f}(y)$ described in Example 2 is called a *kernel density estimator* (KDE). KDEs are biased and, in general, there is no way around this. The introduction of bias is a concession we make in order to make progress on our problem. All interesting problems require some form of concession; a common concession is to assume that the data arise from a *parametric family*.

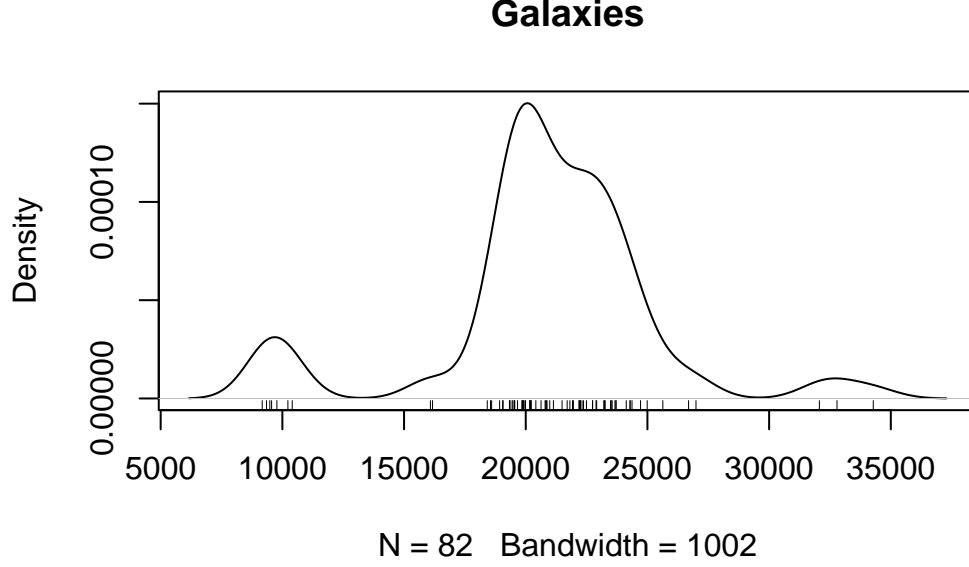


Figure 1: Density estimate for the `galaxy` dataset.

Example 3: Linear Regression

Suppose that $\mathcal{D} = \{(x_1, Y_1), \dots, (x_N, Y_N)\}$ with the x_i 's being fixed vectors in \mathbb{R}^P , with G satisfying the restrictions

$$Y_i = r(x_i) + \epsilon_i, \quad \text{Var}(\epsilon_i) = \sigma^2 < \infty.$$

Without making further assumptions, this sort of problem arises frequently in various *machine learning* problems, with the goal of estimating the mean function $r(x)$ to provide predictions on new data. As statisticians, we often make the further assumption that

$$r(x_i) = x_i^\top \beta, \tag{1}$$

for some *parameter vector* β . If this assumption fails, we will incur bias in estimation of $r(x)$, among other potential problems. We might assume a model like (1) for many reasons:

1. The model (1) is *interpretable*, as by this point you will all be familiar with how to interpret the coefficients of a linear model.
2. The usual least-squares estimator $\hat{\beta} = \arg \min_{\beta} \|\mathbf{Y} - \mathbf{X}\beta\|^2 = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y}$ can be computed efficiently and is relatively stable.

We describe three types of models we might adopt:

- A **parametric model** is a set of distributions $\mathcal{G} = \{g_\theta : \theta \in \Theta\}$ where $\theta \in \mathbb{R}^P$ for some P (technically, we also should assume that the map $\theta \mapsto g_\theta$ is “nice” in some sense, e.g., continuous). We will spend a lot of time with parametric models, particularly generalized linear models (GLMs), mainly because I think it would be irresponsible for you to get through a PhD program without knowing a lot about GLMs (and not because I particularly like parametric models).
- **Nonparametric models** have many different definitions depending on the context one is operating in, but a textbook definition is that a nonparametric model is a set of distributions \mathcal{G} such that any distribution G can be approximated arbitrarily well by a member of \mathcal{G} . For example, the set of all distributions associated with a continuous density $f(x)$ is a nonparametric model (note that even discrete distributions can be well-approximated by continuous distributions with very spiky densities). This is the definition that I’ll adopt for the purposes of this class. We will spend a decent amount of time on nonparametric models near the end of the class.
- A **semiparametric model** is a model that is neither parametric nor nonparametric. Usually, we can write $\mathcal{G} = \{g_{\theta,\eta} : \theta \in \Theta, \eta \in H\}$ where $\theta \in \mathbb{R}^P$ is finite-dimensional while η is an infinite-dimensional parameter. For example, the linear regression model in Example 3 is a semiparametric model, as we have not imposed any restrictions on the distribution of ϵ_i aside from the assumption that it has mean zero and finite variance, and the infinite dimensional parameter η is the conditional density $f(\epsilon | x)$ that is restricted to have mean 0 and variance σ^2 for all x . Semiparametric models try to get the benefits of interpretability of parametric models while not being as susceptible to model misspecification. We will spend a lot of time with semiparametric models, as I think they can usually be used “for free” in place of parametric models.

4 Uncertainty Quantification

The “first-order” job of a statistician is make predictions and help aide in making informed decisions using data: for example, I might predict ice cream sales of 300 cones on a particular day, based on historical data, and recommend buying a certain amount of materials accordingly. A “second-order” job is to further make sure that their predictions/recommendations are robust in some sense; a common way to do this is to attach error bars to the predictions we make (e.g., “I predict 300 cones, but it should be somewhere between 250 and 350”). This is referred to as *uncertainty quantification*. Such UQ problems could refer to single isolated predictions (where we use prediction intervals), or could be predictions about what one would see in a hypothetical infinite sample (where we use confidence/credible intervals).

The philosophical question of how to quantify our “uncertainty” in our conclusions has been widely debated. The two most popular approaches are to quantify uncertainty through the sampling distribution of the data \mathcal{D} (*Frequentist inference*) or through a posterior distribution for the parameter θ (*Bayesian inference*).

We won't spend much time arguing for one approach over the other. My personal belief is that debating the *merits* of the two approaches is largely a distraction, and that it is a bad sign if any methodology you want to use depends fundamentally on philosophical considerations. On the other hand, I think that the two methods often can *complement* each other, as considering a problem from both perspectives can lead to a better overall understanding of that problem. We will see this, in particular, when we consider the problems of multiple testing, normal means estimation, and high-dimensional generalized linear models.

Frequentist Uncertainty Quantification

The Frequentist approach makes use of the sampling distribution $\{G_\theta : \theta \in \Theta\}$ to perform inference. Frequentist methodology attempts to make guarantees about methods in terms of *repeated experiments* — if we were to repeat *exactly the same* experiment $\mathcal{D} \sim G_{\theta_0}$, can we create methods which are guaranteed to perform well even if θ_0 is unknown?

For example, we might aim to construct an interval $[L(\mathcal{D}), U(\mathcal{D})] = [L, U]$ such that, for some parameter of interest θ_j , $L \leq \theta_j \leq U$ holds with some specified probability $1 - \alpha$. Ideally, we would like to choose (L, U) so that this holds *irrespective of the true value of θ* , i.e.,

$$\inf_{\theta \in \Theta} G_\theta(L \leq \theta_j \leq U) = 1 - \alpha. \quad (2)$$

That is, no matter which θ we take, we are guaranteed that our interval covers with probability at least $1 - \alpha$. Often, this goal is a bit too ambitious, and instead we ask only that (2) holds asymptotically with respect to the size of the data N on bounded subsets of Θ , i.e., we ask that $\inf_{\theta \in K} G_\theta(L \leq \theta_j \leq U) = 1 - \alpha + o(1)$ holds for all compact sets K . Fundamental to the Frequentist paradigm is that the methods behave well *uniformly* in θ to the extent possible, in order to account for the fact that θ is unknown.

Bayesian Inference for Uncertainty Quantification

The *Bayesian* approach to probabilistic modeling, by contrast, specifies a *prior* distribution Π on the data generating process G . This typically occurs by way of a *prior density* $\pi(\theta)$ on a *parametric family* $\{G_\theta : \theta \in \Theta\}$ where Θ is a subset of \mathbb{R}^P .

We then apply Bayes rule to obtain the posterior distribution:

$$\pi(\theta \mid \mathcal{D}) = \frac{\pi(\theta) L(\theta)}{m_\pi(\mathcal{D})} \quad \text{where} \quad m_\pi(\mathcal{D}) = \int \pi(\theta) L(\theta) d\theta,$$

and $L(\theta)$ (which tacitly depends on \mathcal{D}) denotes the likelihood function of θ ($L(\theta) = g_\theta(\mathcal{D})$). The posterior distribution $\pi(\theta \mid \mathcal{D})$ can then be used to quantify our uncertainty in θ in terms of probabilities.

There are many ways that folks have tried to make sense of what the posterior probabilities represent philosophically. I endorse the following interpretation:

Claim: The posterior distribution $\pi(\theta \mid \mathcal{D})$ describes what a perfectly-rational robot would believe about θ if (i) the prior $\pi(\theta)$ described their subjective beliefs about θ prior to observing data, and (ii) the only thing they knew about the external world was that $\mathcal{D} \sim G_\theta$ for some $\theta \in \Theta$ (and they knew this with 100% certainty).

Not everyone will agree with this interpretation, but I think it has some features that make it useful to anchor our understanding to. It suggests that we should not interpret posteriors as *our* rational beliefs about θ , but rather the beliefs of a particular, perfectly rational, agent. It also gives us avenues for model criticism, in two ways: we can criticize the choice of the prior $\pi(\theta)$ in (i), or we can criticize the choice of G_θ in (ii). It also reminds us that the output of Bayesian models themselves are operating under very strong assumptions: *our robot believes the model with 100% certainty, and so can afford to behave in ways that we might deem irrational to someone who recognizes that this is not the case.*

A personal gripe I have with Bayesian procedures is that, due to their overconfidence in the correctness of the likelihood, the uncertainty quantification we get may not be very robust - it is often the case that point estimates from Bayesian methods will be pretty good even when the model is wrong, but the uncertainty quantification can be poor (either in Frequentist terms, or relative to what you would get if you had used a correct likelihood).

5 Computation via Markov chain Monte Carlo

You will be exposed to Bayesian computation in other courses. On the off chance that you have not covered this material yet, I will review the high-level idea of Markov chain Monte Carlo.

Our ultimate goal is to obtain inferences for θ based on the posterior distribution $\pi(\theta \mid \mathcal{D})$. We might be interested, for example, in the *Bayes estimator* for θ , given by

$$\tilde{\theta} = \mathbb{E}_\pi(\theta \mid \mathcal{D}) = \int \theta \pi(\theta \mid \mathcal{D}) d\theta.$$

The catch is that integrals like this are often computationally intractable. If we could generate a *sample* $\theta_1, \dots, \theta_B \stackrel{\text{iid}}{\sim} \pi(\theta \mid \mathcal{D})$ from the posterior, however, then we could approximate this expectation as $\tilde{\theta} \approx B^{-1} \sum_{b=1}^B \theta_b$. We could also approximate a $100(1 - 2\alpha)\%$ credible interval for θ by taking the α^{th} and $(1 - \alpha)^{\text{th}}$ sample quantiles of the θ_b 's. These are just examples; we can basically compute whatever features of $\pi(\theta \mid \mathcal{D})$ we want if we have a sample from the posterior.

Unfortunately, sampling from $\pi(\theta \mid \mathcal{D})$ is (in general) no easier than computing integrals. The idea behind Markov chain Monte Carlo (MCMC) is to replace the samples $\theta_1, \dots, \theta_B \stackrel{\text{iid}}{\sim} \pi(\theta \mid \mathcal{D})$

with a *Markov chain* such that $\theta_b \sim q(\theta \mid \theta_{b-1})$. The distribution $q(\theta \mid \theta')$ is called a *Markov transition function* (MTF), and as long as the MTF leaves the posterior invariant

$$\pi(\theta \mid \mathcal{D}) = \int q(\theta \mid \theta') \pi(\theta' \mid \mathcal{D}) d\theta'$$

and satisfies some other extremely minor technical conditions, the samples $\theta_1, \dots, \theta_B$ will function more-or-less like a sample from the posterior. There are two catches.

1. The samples are no longer independent, so we may have to take a (much) larger B to get reasonable approximations.
2. The samples are no longer distributed exactly according to $\pi(\theta \mid \mathcal{D})$.

Both of these issues are related to how fast the chain *mixes*, i.e., how quickly the chain “forgets” its history.

To address the second issue, we typically *burn in* the chain by discarding (say) the first 1000 samples from the chain, the idea being that we should be pretty close to $\pi(\theta \mid \mathcal{D})$ at that point. The number 1000 I just mentioned is arbitrary, and the correct burn in sample size can range from less than 10 (for good chains) to larger than the number of particles in the observable universe (for slow mixing chains, and no, I am not exaggerating).

Upon reflection, the first issue is not too different from the first, but it is generally resolved in a different way. One approach is to *thin* the chain, retaining only (say) every 10th sample, and then treat the samples as approximately independent. Again, 10 is arbitrary. My personal opinion is that thinning is a waste of time *unless you are running out of RAM*. The better solution is to explicitly account for dependence in the samples in your assessments of your effective sample size, which will typically be returned by whatever software you are using.

6 Exercises

Misc

Exercise 1: Density Estimation

Consider the density estimation problem with the **galaxy** dataset. Suppose your goal is to estimate the number of modes of the density.

- a. How would you go about estimating the number of modes? (Note: you might use the KDE in some fashion, or you might suggest some alternative ideas; no wrong answers here.)
- b. Try out the KDE with different choices of h (consider both very big and very small values of h). How does h relate to the number of modes you end up recovering?

Does it seem important?

- c. Without actually doing so, give some ideas for how you might go about assessing how well your estimator from (a) can be expected to perform. Bonus points if you can come up with a method for quantifying uncertainty.

Review of Bayesian Inference in Simple Conjugate Families

Exercise 2: Beta-Binomial

Suppose X_1, \dots, X_N are iid Bernoulli random variables with success probability p (i.e., the X_i 's are the result of flipping a biased coin with probability of heads p). Suppose that p is given a $B(\alpha, \beta)$ prior distribution, having density

$$\pi(p) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} p^{\alpha-1} (1-p)^{\beta-1} I(0 \leq p \leq 1).$$

Derive the posterior of $[p \mid X_1, \dots, X_N]$.

Exercise 3: Dirichlet-Multinomial

Suppose X_1, \dots, X_N are iid categorical random variables taking values in $\{1, \dots, K\}$ with probabilities $p = (p_1, \dots, p_K)$ respectively; the likelihood of this model is $p \sim \text{Dirichlet}(\alpha_1, \dots, \alpha_K)$ which has density

$$\pi(p) = \frac{\Gamma(\sum_k \alpha_k)}{\prod_k \Gamma(\alpha_k)} \prod_{k=1}^K p_k^{\alpha_k-1},$$

on the *simplex* $\mathbb{S}_{K-1} = \{p : p_k \geq 0, \sum_k p_k = 1\}$; this is a density on \mathbb{R}^{K-1} with $p_K \equiv 1 - \sum_{k=1}^{K-1} p_k$. Find the posterior distribution of $[p \mid X_1, \dots, X_N]$ (it may be helpful to define $n_k = \sum_{i=1}^N I(X_i = k)$).

Exercise 4: Properties of the Gamma Distribution

We say that X has a *gamma* distribution if X has density

$$\frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x} I(x > 0),$$

and we write $X \sim \text{Gam}(\alpha, \beta)$. Suppose that $X \sim \text{Gam}(\alpha, \beta)$ and $Y \sim \text{Gam}(\beta, \beta)$ and that X is independent of Y .

- a. Let $W = X + Y$ and $Z = X/(X + Y)$. Show that W and Z are independent with

$W \sim \text{Gam}(\alpha + \beta, b)$ and $Z \sim \text{B}(\alpha, \beta)$.

- b. Suppose that we have access to a random number generator (RNG) capable of producing independent $\text{Gam}(a, b)$ random variables (such as the `rgamma` function in R) for any choice of a and b . Explain how to use this RNG to sample $\text{B}(\alpha, \beta)$ random variables.

Exercise 5: Normal-Normal

Suppose $X_1, \dots, X_N \stackrel{\text{iid}}{\sim} \text{Normal}(\theta, \sigma_0^2)$ where σ^2 is *known*. Suppose that θ is given a normal prior distribution with mean m and variance v . Derive the posterior distribution of $[\theta \mid X_1, \dots, X_N]$.

Exercise 6: Normal-Gamma

Suppose $X_1, \dots, X_N \stackrel{\text{iid}}{\sim} \text{Normal}(\theta, \sigma^2)$ with θ *known* but σ^2 *unknown*. Suppose that $\omega = \sigma^{-2}$ has a $\text{Gam}(\alpha, \beta)$ prior. Derive the posterior distribution of $[\omega \mid X_1, \dots, X_N]$.

7 Some Comments on Notation

I have a (bad) habit of using notation without considering that students may not be aware of some of it. For your benefit, I'll give some of the usual notation that I might assume you know. It is standard notation that you are likely to see in papers, but maybe unlikely to have seen prior to this point.

None of this really matters for the purpose of this course, but it is easier for me to just tell you what the notation means than stop myself from using it when I feel like it.

- If X and Y are random variables depending on an index N (often the sample size) then the statement $Y = o_P(X)$ means that $Y/X \rightarrow 0$ in probability as $N \rightarrow \infty$. For example, the weak law of large numbers can be expressed compactly as

$$\frac{1}{N} \sum_i X_i = \mu + o_P(1) \quad \text{or possibly} \quad \sum_i X_i = N\mu + o_P(N).$$

- The statement $Y = O_P(X)$ means that Y/X is *bounded in probability*. This means that (i) for every $\epsilon > 0$ there (ii) exists a positive constant K such that (iii) for sufficiently large N we have $\Pr(|Y| \leq K|X|) \geq 1 - \epsilon$. An implication of the central limit theorem is that

$$\frac{1}{N} \sum_i X_i = \mu + O_P(N^{-1/2})$$

because $N^{1/2}(\bar{X} - \mu)$ converges in distribution to a normal distribution.

- The theories of discrete and continuous variables are unified by the *measure theoretic approach to probability*, which we don't require you to know. Within this framework, the expected value of a random variable $X \sim F$ is written

$$\mathbb{E}(X) = \int x F(dx).$$

When X is continuous (or discrete) this quantity happens to be equal to

$$\int x f(x) dx \quad \text{or} \quad \sum_x x f(x),$$

where $f(x)$ is the density (or mass) function of X .

Because the discrete and continuous settings are effectively the same, I may write things like $\int x f(x) dx$ even when X is discrete.